



Experiment No. 3
Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model
Date of Performance:07-08-23
Date of Submission:17-08-23

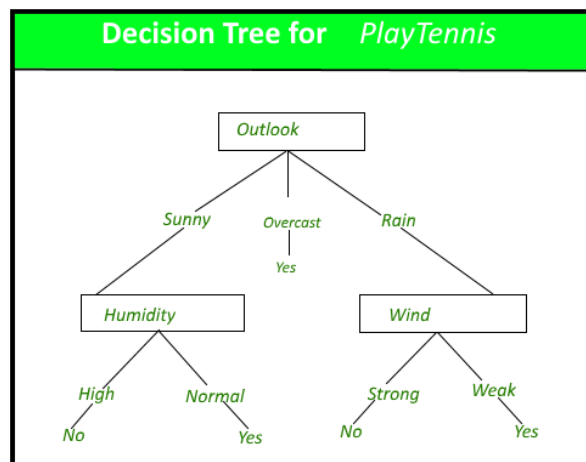


Aim: Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: To perform various feature engineering tasks, apply Decision Tree Algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score. Improve the performance by performing different data engineering and feature engineering tasks.

Theory:

Decision Tree is the most powerful and popular tool for classification and prediction. A Decision tree is a flowchart-like tree structure, where each internal node denotes a test on an attribute, each branch represents an outcome of the test, and each leaf node (terminal node) holds a class label.



Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:



>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras,



Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.

Code:

Conclusion:

The Decision Tree model showed strong performance on the Adult Census Income Dataset. It effectively handled categorical attributes through one-hot encoding and conducted essential data preprocessing tasks, such as managing missing values, eliminating irrelevant columns, and segregating features to enhance the model's effectiveness.

Optimizing hyperparameters is crucial to further boost the Decision Tree model's performance, as it allows us to fine-tune the model's complexity by setting limits on various parameters. To elevate the model's performance, we can fine-tune hyperparameters like max depth, min samples split, etc., using techniques like Grid Search or Random Search.

Regarding model evaluation:

Accuracy: The model achieved an accuracy of 0.85, implying that approximately 85% of its predictions were correct.

Confusion Matrix: The confusion matrix shows that there were 9860 true positive predictions, 481 false positive predictions, 1823 false negatives, and 1646 true negative predictions.

Precision: The precision value of 0.84 suggests that among the instances predicted as class 0, about 84% were correctly predicted as class 0, while approximately 16% were falsely classified as class 1.

Recall: With a recall of 0.95, the model effectively captured a significant portion of instances belonging to class 0, indicating that it correctly identified 95% of them. However, it captured only 47% of instances from class 1.



-F1 Score: The F1 score is a balanced measure that combines precision and recall. The F1 score of 0.90 for class 0 indicates a good balance between precision and recall for this class, while the F1 score of 0.59 for class 1 suggests a moderate balance between precision and recall for this class in the model's performance.

```
# Import libraries
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
# To ignore warning messages
import warnings
warnings.filterwarnings('ignore')
# Adult dataset path
adult_dataset_path = "/content/adult.csv"
```

```
# Function for loading adult dataset
def load_adult_data(adult_path=adult_dataset_path):
    csv_path = os.path.join(adult_path)
    return pd.read_csv(csv_path)
```

```
# Calling load adult function and assigning to a new variable df
df = load_adult_data()
# load top 3 rows values from adult dataset
df.head(3)
```

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relation
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husl

```
print ("Rows : ",df.shape[0])
print ("Columns : ",df.shape[1])
print ("\nFeatures : \n",df.columns.tolist())
print ("\nMissing values : ", df.isnull().sum().values.sum())
print ("\nUnique values : \n",df.nunique())
```

```
Rows : 48842
Columns : 15
```

```
Features :
['age', 'workclass', 'fnlwgt', 'education', 'educational-num', 'marital-status', 'occupation', 'relationship', 'race', 'gender', 'capital-
```

```
Missing values : 0
```

```
Unique values :
age                74
workclass           9
fnlwgt            28523
education          16
educational-num    16
marital-status      7
occupation         15
relationship        6
race                5
gender              2
capital-gain       123
capital-loss        99
hours-per-week     96
native-country     42
income             2
dtype: int64
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48842 entries, 0 to 48841
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   48842 non-null  int64
```

```

1  workclass      48842 non-null object
2  fnlwgt         48842 non-null int64
3  education      48842 non-null object
4  educational-num 48842 non-null int64
5  marital-status 48842 non-null object
6  occupation     48842 non-null object
7  relationship   48842 non-null object
8  race           48842 non-null object
9  gender         48842 non-null object
10 capital-gain   48842 non-null int64
11 capital-loss   48842 non-null int64
12 hours-per-week 48842 non-null int64
13 native-country 48842 non-null object
14 income         48842 non-null object

```

```

dtypes: int64(6), object(9)
memory usage: 5.6+ MB

```

```
df.describe()
```

	age	fnlwgt	educational-num	capital-gain	capital-loss	hours
count	48842.000000	4.884200e+04	48842.000000	48842.000000	48842.000000	48842.0
mean	38.643585	1.896641e+05	10.078089	1079.067626	87.502314	40.4
std	13.710510	1.056040e+05	2.570973	7452.019058	403.004552	12.3
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.0
25%	28.000000	1.175505e+05	9.000000	0.000000	0.000000	40.0
50%	37.000000	1.781445e+05	10.000000	0.000000	0.000000	40.0
75%	48.000000	2.376420e+05	12.000000	0.000000	0.000000	45.0

```
df.head()
```

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relation
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husi
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husi
3	44	Private	160323	Some-	10	Married-civ-	Machine-	Husi

```
# checking "?" total values present in particular 'workclass' feature
```

```
df_check_missing_workclass = (df['workclass']=='?').sum()
```

```
df_check_missing_workclass
```

```
2799
```

```
# checking "?" total values present in particular 'occupation' feature
```

```
df_check_missing_occupation = (df['occupation']=='?').sum()
```

```
df_check_missing_occupation
```

```
2809
```

```
# checking "?" values, how many are there in the whole dataset
```

```
df_missing = (df=='?').sum()
```

```
df_missing
```

```

age           0
workclass     2799
fnlwgt        0
education     0
educational-num 0
marital-status 0
occupation    2809

```

```

relationship    0
race            0
gender          0
capital-gain    0
capital-loss    0
hours-per-week  0
native-country  857
income          0
dtype: int64

```

```

percent_missing = (df=='?').sum() * 100/len(df)
percent_missing

```

```

age            0.000000
workclass      5.730724
fnlwgt         0.000000
education      0.000000
educational-num 0.000000
marital-status 0.000000
occupation     5.751198
relationship    0.000000
race           0.000000
gender         0.000000
capital-gain    0.000000
capital-loss    0.000000
hours-per-week  0.000000
native-country  1.754637
income         0.000000
dtype: float64

```

```

# find total number of rows which doesn't contain any missing value as '?'
df.apply(lambda x: x != '?',axis=1).sum()

```

```

age            48842
workclass      46043
fnlwgt         48842
education      48842
educational-num 48842
marital-status 48842
occupation     46033
relationship    48842
race           48842
gender         48842
capital-gain    48842
capital-loss    48842
hours-per-week  48842
native-country  47985
income         48842
dtype: int64

```

```

# dropping the rows having missing values in workclass
df = df[df['workclass'] != '?']
df.head()

```

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relation:
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husl
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husl
3	44	Private	160323	Some-	10	Married-civ-	Machine-	Husl

```

# select all categorical variables
df_categorical = df.select_dtypes(include=['object'])
# checking whether any other column contains '?' value
df_categorical.apply(lambda x: x=='?',axis=1).sum()

```

```

workclass      0
education      0
marital-status  0

```



```

occupation    10
relationship  0
race          0
gender        0
native-country 811
income        0
dtype: int64

```

```

from sklearn import preprocessing
# encode categorical variables using Label Encoder
# select all categorical variables
df_categorical = df.select_dtypes(include=['object'])
df_categorical.head()

```

	workclass	education	marital-status	occupation	relationship	race	gender	native-country
0	Private	11th	Never-married	Machine-op-inspct	Own-child	Black	Male	United-States
1	Private	HS-grad	Married-civ-spouse	Farming-fishing	Husband	White	Male	United-States
			Married-					

```

# apply label encoder to df_categorical
le = preprocessing.LabelEncoder()
df_categorical = df_categorical.apply(le.fit_transform)
df_categorical.head()

```

	workclass	education	marital-status	occupation	relationship	race	gender	native-country
0	2	1	4	6	3	2	1	39
1	2	11	2	4	0	4	1	39
2	1	7	2	10	0	4	1	39
3	2	15	2	6	0	2	1	39

```

# Next, Concatenate df_categorical dataframe with original df (dataframe)
# first, Drop earlier duplicate columns which had categorical values
df = df.drop(df_categorical.columns,axis=1)
df = pd.concat([df,df_categorical],axis=1)
df.head()

```

	age	fnlwgt	educational-num	capital-gain	capital-loss	hours-per-week	workclass	education	marital-status
0	25	226802	7	0	0	40	2	1	
1	38	89814	9	0	0	50	2	11	
2	28	336951	12	0	0	40	1	7	
3	44	160323	10	7688	0	40	2	15	

```

# look at column type
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 46033 entries, 0 to 48841
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   46033 non-null  int64
1   fnlwgt                46033 non-null  int64
2   educational-num       46033 non-null  int64
3   capital-gain          46033 non-null  int64
4   capital-loss          46033 non-null  int64
5   hours-per-week        46033 non-null  int64
6   workclass             46033 non-null  int64
7   education             46033 non-null  int64
8   marital-status        46033 non-null  int64

```

```

9  occupation      46033 non-null  int64
10 relationship    46033 non-null  int64
11 race            46033 non-null  int64
12 gender          46033 non-null  int64
13 native-country  46033 non-null  int64
14 income          46033 non-null  int64
dtypes: int64(15)
memory usage: 5.6 MB

```

```

# convert target variable income to categorical
df['income'] = df['income'].astype('category')
# check df.info again whether everything is in right format or not
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 46033 entries, 0 to 48841
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    46033 non-null  int64
1   fnlwgt                 46033 non-null  int64
2   educational-num        46033 non-null  int64
3   capital-gain           46033 non-null  int64
4   capital-loss           46033 non-null  int64
5   hours-per-week         46033 non-null  int64
6   workclass              46033 non-null  int64
7   education              46033 non-null  int64
8   marital-status         46033 non-null  int64
9   occupation             46033 non-null  int64
10  relationship           46033 non-null  int64
11  race                   46033 non-null  int64
12  gender                 46033 non-null  int64
13  native-country         46033 non-null  int64
14  income                 46033 non-null  category
dtypes: category(1), int64(14)
memory usage: 5.3 MB

```

```

# Importing train_test_split
from sklearn.model_selection import train_test_split
# Putting independent variables/features to X
X = df.drop('income',axis=1)
# Putting response/dependent variable/feature to y
y = df['income']

```

```
X.head(3)
```

	age	fnlwgt	educational-num	capital-gain	capital-loss	hours-per-week	workclass	education	marital-status
0	25	226802	7	0	0	40	2	1	
1	38	89814	9	0	0	50	2	11	

```
y.head(3)
```

```

0    0
1    0
2    1
Name: income, dtype: category
Categories (2, int64): [0, 1]

```

```

# Splitting the data into train and test
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.30,random_state=99)
X_train.head()

```

```

hours-
# Importing decision tree classifier from sklearn library
from sklearn.tree import DecisionTreeClassifier
# Fitting the decision tree with default hyperparameters, apart from
# max_depth which is 5 so that we can plot and read the tree.
dt_default = DecisionTreeClassifier(max_depth=5)
dt_default.fit(X_train,y_train)

```

```

DecisionTreeClassifier
DecisionTreeClassifier(max_depth=5)

```

```

# check the evaluation metrics of our default model
# Importing classification report and confusion matrix from sklearn metrics
from sklearn.metrics import classification_report,confusion_matrix,accuracy_score
# making predictions
y_pred_default = dt_default.predict(X_test)
# Printing classifier report after prediction
print(classification_report(y_test,y_pred_default))

```

	precision	recall	f1-score	support
0	0.86	0.95	0.90	10341
1	0.79	0.53	0.64	3469
accuracy			0.85	13810
macro avg	0.83	0.74	0.77	13810
weighted avg	0.84	0.85	0.84	13810

```

# Printing confusion matrix and accuracy
print(confusion_matrix(y_test,y_pred_default))
print(accuracy_score(y_test,y_pred_default))

```

```

[[9861 480]
 [1616 1853]]
0.848225923244026

```

```
!pip install my-package
```

```

Collecting my-package
  Downloading my_package-0.0.0-py3-none-any.whl (2.0 kB)
Installing collected packages: my-package
Successfully installed my-package-0.0.0

```

```
!pip install pydotplus
```

```

Requirement already satisfied: pydotplus in /usr/local/lib/python3.10/dist-packages (2.0.2)
Requirement already satisfied: pyparsing>=2.0.1 in /usr/local/lib/python3.10/dist-packages (from pydotplus) (3.1.1)

```

```

# Importing required packages for visualization
from IPython.display import Image
from six import StringIO
from sklearn.tree import export_graphviz
import pydotplus,graphviz
# Putting features
features = list(df.columns[1:])
features

```

```


['fnlwgt',
 'educational-num',
 'capital-gain',
 'capital-loss',
 'hours-per-week',
 'workclass',
 'education',
 'marital-status',
 'occupation',
 'relationship',
 'race',
 'gender',
 'native-country',
 'income']

```


```
!pip install graphviz
```

Requirement already satisfied: graphviz in /usr/local/lib/python3.10/dist-packages (0.20.1)

```
# plotting tree with max_depth=3
dot_data = StringIO()
export_graphviz(dt_default, out_file=dot_data,
feature_names=features, filled=True, rounded=True)
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
Image(graph.create_png())
```



```
# GridSearchCV to find optimal max_depth
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV
# specify number of folds for k-fold CV
n_folds = 5
# parameters to build the model on
parameters = {'max_depth': range(1, 40)}
# instantiate the model
dtree = DecisionTreeClassifier(criterion = "gini",
random_state = 100)
# fit tree on training data
tree = GridSearchCV(dtree, parameters,
cv=n_folds,
scoring="accuracy")
tree.fit(X_train, y_train)
```




```
# scores of GridSearch CV
scores = tree.cv_results_
pd.DataFrame(scores).head()
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_max_depth
0	0.017508	0.001094	0.003653	0.000255	1
1	0.026208	0.000638	0.003497	0.000050	2
2	0.035516	0.000600	0.003571	0.000109	3
3	0.046339	0.002691	0.003698	0.000265	4
4	0.054046	0.000881	0.003738	0.000126	5

```
# GridSearchCV to find optimal max_depth
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV
# specify number of folds for k-fold CV
n_folds = 5
# parameters to build the model on
parameters = {'min_samples_leaf': range(5, 200, 20)}
# instantiate the model
```


```
dtree = DecisionTreeClassifier(criterion = "gini",
random_state = 100)
# fit tree on training data
tree = GridSearchCV(dtree, parameters,
cv=n_folds,
scoring="accuracy")
tree.fit(X_train, y_train)
```



```
# scores of GridSearch CV
scores = tree.cv_results_
pd.DataFrame(scores).head()
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_min_samples
0	0.137336	0.004291	0.004951	0.000940	
1	0.115085	0.003973	0.004306	0.000304	
2	0.107868	0.005090	0.004147	0.000113	
3	0.101131	0.002170	0.004166	0.000221	
4	0.100600	0.004742	0.004072	0.000051	

```
# GridSearchCV to find optimal min_samples_split
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV
# specify number of folds for k-fold CV
n_folds = 5
# parameters to build the model on
parameters = {'min_samples_split': range(5, 200, 20)}
# instantiate the model
dtree = DecisionTreeClassifier(criterion = "gini",
random_state = 100)
# fit tree on training data
tree = GridSearchCV(dtree, parameters,
cv=n_folds,
scoring="accuracy")
tree.fit(X_train, y_train)
```



```
# scores of GridSearch CV
scores = tree.cv_results_
pd.DataFrame(scores).head()
```

mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_min_samples
---------------	--------------	-----------------	----------------	-------------------

```
# Create the parameter grid
param_grid = {
    'max_depth': range(5, 15, 5),
    'min_samples_leaf': range(50, 150, 50),
    'min_samples_split': range(50, 150, 50),
    'criterion': ["entropy", "gini"]
}
n_folds = 5

# Instantiate the grid search model
dtree = DecisionTreeClassifier()
grid_search = GridSearchCV(estimator = dtree, param_grid = param_grid,
cv = n_folds, verbose = 1)
# Fit the grid search to the data
grid_search.fit(X_train,y_train)
```

Fitting 5 folds for each of 16 candidates, totalling 80 fits

```
GridSearchCV
  estimator: DecisionTreeClassifier
    DecisionTreeClassifier
```

```
# cv results
cv_results = pd.DataFrame(grid_search.cv_results_)
cv_results
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_criterion
0	0.059965	0.001689	0.004816	0.000741	entropy
1	0.058885	0.001184	0.004362	0.000257	entropy
2	0.059616	0.004424	0.003817	0.000191	entropy
3	0.056671	0.000835	0.003579	0.000023	entropy
4	0.095637	0.003662	0.003889	0.000054	entropy
5	0.127947	0.018810	0.005677	0.000847	entropy
6	0.137194	0.005227	0.006264	0.001309	entropy
7	0.122113	0.018699	0.004998	0.000869	entropy

```
# printing the optimal accuracy score and hyperparameters
```

```
print("best accuracy", grid_search.best_score_)
```

```
print(grid_search.best_estimator_)
```

```
best accuracy 0.8523105983446813
```

```
DecisionTreeClassifier(max_depth=10, min_samples_leaf=50, min_samples_split=50)
```

```
# model with optimal hyperparameters
```

```
clf_gini = DecisionTreeClassifier(criterion = "gini",
```

```
random_state = 100,
```

```
max_depth=10,
```

```
min_samples_leaf=50,
```

```
min_samples_split=50)
```

```
clf_gini.fit(X_train, y_train)
```

```
DecisionTreeClassifier
DecisionTreeClassifier(max_depth=10, min_samples_leaf=50, min_samples_split=50,
random_state=100)
```

```
# accuracy score
```

```
clf_gini.score(X_test,y_test)
```

```
0.852860246198407
```

```
#plotting the tree
```

```
dot_data = StringIO()
```

```
export_graphviz(clf_gini, out_file=dot_data,feature_names=features,filled=True,rounded=True)
```

```
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
```

```
Image(graph.create_png())
```

```
# tree with max_depth = 3
```

```
clf_gini = DecisionTreeClassifier(criterion = "gini",
```

```
random_state = 100,
```

```
max_depth=3,
```

```
min_samples_leaf=50,
```

```
min_samples_split=50)
```

```

clf_gini.fit(X_train, y_train)
# score
print(clf_gini.score(X_test,y_test))
# plotting tree with max_depth=3
dot_data = StringIO()
export_graphviz(clf_gini, out_file=dot_data,feature_names=features,filled=True,rounded=True)


```

0.8331643736422882

```

# classification metrics
from sklearn.metrics import classification_report,confusion_matrix
y_pred = clf_gini.predict(X_test)
print(classification_report(y_test, y_pred))
# confusion matrix
print(confusion_matrix(y_test,y_pred))

```



	precision	recall	f1-score	support
0	0.84	0.95	0.90	10341
1	0.77	0.47	0.59	3469
accuracy			0.83	13810
macro avg	0.81	0.71	0.74	13810
weighted avg	0.83	0.83	0.82	13810

```

[[9860 481]
 [1823 1646]]

```