



Experiment No. 6
Apply Boosting Algorithm on Adult Census Income Dataset and analyze the performance of the model
Date of Performance:04-09-23
Date of Submission:25-09-23



Aim: Apply Boosting algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: Apply Boosting algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score.

Theory:

Suppose that as a patient, you have certain symptoms. Instead of consulting one doctor, you choose to consult several. Suppose you assign weights to the value or worth of each doctor's diagnosis, based on the accuracies of previous diagnosis they have made. The final diagnosis is then a combination of the weighted diagnosis. This is the essence behind boosting.

Algorithm: Adaboost- A boosting algorithm—create an ensemble of classifiers. Each one gives a weighted vote.

Input:

- D , a set of d class labelled training tuples
- k , the number of rounds (one classifier is generated per round)
- a classification learning scheme

Output: A composite model

Method

1. Initialize the weight of each tuple in D is $1/d$
2. For $i=1$ to k do // for each round
3. Sample D with replacement according to the tuple weights to obtain D_i
4. Use training set D_i to derive a model M_i
5. Compute $\text{error}(M_i)$, the error rate of M_i
6. $\text{Error}(M_i) = \sum w_j \cdot \text{err}(X_j)$
7. If $\text{Error}(M_i) > 0.5$ then
8. Go back to step 3 and try again
9. endif
10. for each tuple in D_i that was correctly classified do
11. Multiply the weight of the tuple by $\text{error}(M_i)/(1-\text{error}(M_i))$
12. Normalize the weight of each tuple
13. end for

To use the ensemble to classify tuple X

1. Initialize the weight of each class to 0



2. for $i=1$ to k do // for each classifier
3. $w_i = \log((1 - \text{error}(M_i)) / \text{error}(M_i))$ // weight of the classifiers vote
4. $C = M_i(X)$ // get class prediction for X from M_i
5. Add w_i to weight for class C
6. end for
7. Return the class with the largest weight.

Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.



hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad & Tobago, Peru, Hong, Holand-Netherlands.

Code:

Conclusion:

The accuracy of the model stands at 0.865, indicating that it effectively predicts income levels. The confusion matrix reveals that there are 637 true positive predictions, 144 true negatives, 379 false positives, and 854 false negatives.

In terms of precision, the model achieves a precision of 0.88 for predicting income level 0 and 0.79 for predicting income level 1. For recall, it attains a recall of 0.94 for income level 0 and 0.63 for income level 1. The F1-score, which combines precision and recall, is 0.91 for income level 0 and 0.70 for income level 1, indicating the model's overall effectiveness.

Both Random Forest and AdaBoost are capable of achieving high accuracy and are less prone to overfitting. However, Random Forest tends to be more robust to changes in hyperparameter tuning and often requires fewer adjustments in this regard. Additionally, Random Forest offers the benefit of feature importance analysis, making it more interpretable. In contrast, AdaBoost's sequential approach may result in lower interpretability.

In scenarios involving imbalanced data, AdaBoost outperforms Random Forest by giving more weight to minority class samples, effectively addressing class imbalance. In summary, AdaBoost and Random Forest are powerful ensemble algorithms, but their performance can vary based on hyperparameter settings and dataset characteristics.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import io
from sklearn.metrics import accuracy_score, precision_score, f1_score, confusion_matrix, classification_report
from sklearn.model_selection import cross_val_score
from sklearn.metrics import mean_squared_error
```

```
import os
for dirname, _, filenames in os.walk('/content/adult.csv'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

```
file = ('/content/adult.csv')
df = pd.read_csv(file)
```

```
print(df.head())
```

```

age workclass  fnlwgt    education  education.num  marital.status \
0    90      ?    77053      HS-grad             9      Widowed
1    82  Private  132870      HS-grad             9      Widowed
2    66      ?  186061  Some-college            10      Widowed
3    54  Private  140359      7th-8th             4      Divorced
4    41  Private  264663  Some-college            10      Separated

      occupation  relationship    race    sex  capital.gain \
0              ?  Not-in-family  White  Female             0
1  Exec-managerial  Not-in-family  White  Female             0
2              ?    Unmarried  Black  Female             0
3  Machine-op-inspct    Unmarried  White  Female             0
4  Prof-specialty    Own-child   White  Female             0

capital.loss  hours.per.week  native.country  income
0          4356             40  United-States  <=50K
1          4356             18  United-States  <=50K
2          4356             40  United-States  <=50K
3          3900             40  United-States  <=50K
4          3900             40  United-States  <=50K
```

```
print(df.info())
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             32561 non-null  object
2   fnlwgt                32561 non-null  int64
3   education              32561 non-null  object
4   education.num          32561 non-null  int64
5   marital.status        32561 non-null  object
6   occupation             32561 non-null  object
7   relationship           32561 non-null  object
8   race                  32561 non-null  object
9   sex                   32561 non-null  object
10  capital.gain           32561 non-null  int64
11  capital.loss           32561 non-null  int64
12  hours.per.week         32561 non-null  int64
13  native.country         32561 non-null  object
14  income                 32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
None
```

```

for i in df.columns:
    t = df[i].value_counts()
    index = list(t.index)
    print ("Count of ? in", i)
    for i in index:
        temp = 0
        if i == '?':
            print (t['?'])
            temp = 1
            break
```

```
if temp == 0:
    print ("0")
```

Streaming output truncated to the last 5000 lines.

[illegible]

```
df=df.loc[(df['workclass'] != '?') & (df['native.country'] != '?')]  
print(df.head())
```

	age	workclass	fnlwgt	education	education.num	marital.status	\
1	82	Private	132870	HS-grad	9	Widowed	
3	54	Private	140359	7th-8th	4	Divorced	
4	41	Private	264663	Some-college	10	Separated	
5	34	Private	216864	HS-grad	9	Divorced	
6	38	Private	150601	10th	6	Separated	

	occupation	relationship	race	sex	capital.gain	\
1	Exec-managerial	Not-in-family	White	Female	0	
3	Machine-op-inspct	Unmarried	White	Female	0	
4	Prof-specialty	Own-child	White	Female	0	
5	Other-service	Unmarried	White	Female	0	
6	Adm-clerical	Unmarried	White	Male	0	

	capital.loss	hours.per.week	native.country	income
1	4356	18	United-States	<=50K
3	3900	40	United-States	<=50K
4	3900	40	United-States	<=50K

```

5      3770      45 United-States <=50K
6      3770      40 United-States <=50K

```

```

df["income"] = [1 if i=='>50K' else 0 for i in df["income"]]
print(df.head())

```

```

   age workclass  fnlwgt  education  education.num marital.status \
1   82   Private  132870    HS-grad           9      Widowed
3   54   Private  140359    7th-8th           4      Divorced
4   41   Private  264663  Some-college        10      Separated
5   34   Private  216864    HS-grad           9      Divorced
6   38   Private  150601     10th           6      Separated

```

```

   occupation  relationship  race  sex  capital.gain \
1  Exec-managerial  Not-in-family  White  Female      0
3  Machine-op-inspct  Unmarried  White  Female      0
4   Prof-specialty  Own-child  White  Female      0
5   Other-service  Unmarried  White  Female      0
6    Adm-clerical  Unmarried  White  Male      0

```

```

   capital.loss  hours.per.week  native.country  income
1         4356           18  United-States      0
3         3900           40  United-States      0
4         3900           40  United-States      0
5         3770           45  United-States      0
6         3770           40  United-States      0

```

```

df_more=df.loc[df['income'] == 1]
print(df_more.head())

```

```

   age  workclass  fnlwgt  education  education.num marital.status \
7   74  State-gov  88638  Doctorate           16  Never-married
10  45   Private  172274  Doctorate           16      Divorced
11  38  Self-emp-not-inc  164526  Prof-school        15  Never-married
12  52   Private  129177  Bachelors            13      Widowed
13  32   Private  136204  Masters             14      Separated

```

```

   occupation  relationship  race  sex  capital.gain \
7  Prof-specialty  Other-relative  White  Female      0
10 Prof-specialty  Unmarried  Black  Female      0
11 Prof-specialty  Not-in-family  White  Male      0
12  Other-service  Not-in-family  White  Female      0
13 Exec-managerial  Not-in-family  White  Male      0

```

```

   capital.loss  hours.per.week  native.country  income
7         3683           20  United-States      1
10         3004           35  United-States      1
11         2824           45  United-States      1
12         2824           20  United-States      1
13         2824           55  United-States      1

```

```

workclass_types = df_more['workclass'].value_counts()
labels = list(workclass_types.index)
aggregate = list(workclass_types)
print(workclass_types)
print(aggregate)
print(labels)

```

```


Private      4876
Self-emp-not-inc  714
Local-gov    609
Self-emp-inc  600
Federal-gov  365
State-gov    344
Name: workclass, dtype: int64
[4876, 714, 609, 600, 365, 344]
['Private', 'Self-emp-not-inc', 'Local-gov', 'Self-emp-inc', 'Federal-gov', 'State-gov']

```


```

plt.figure(figsize=(7,7))
plt.pie(aggregate, labels=labels, autopct='%1.1f%%', shadow = True)
plt.axis('equal')
plt.show()

```



```
#Count plot on single categorical variable
sns.countplot(x='income', data = df)
plt.show()
df['income'].value_counts()
```




```
0    22661
1     7508
Name: income, dtype: int64
```

```
#Plot figsize
plt.figure(figsize=(10,7))
sns.heatmap(df.corr(), cmap='coolwarm', annot=True)
print(plt.show())
plt.figure(figsize=(10,7))
sns.distplot(df['age'], color="red", bins=100)
plt.ylabel("Distribution", fontsize = 10)
plt.xlabel("Age", fontsize = 10)
plt.show()
```



```
<ipython-input-14-c01c35a847eb>:3: FutureWarning: The default value of numeric_c
sns.heatmap(df.corr(), cmap='coolwarm', annot=True)
```



None

```
<ipython-input-14-c01c35a847eb>:6: UserWarning:
```


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see

<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['age'], color="red", bins=100)
```



#To find distribution of categorical columns w.r.t income

```
fig, axes = plt.subplots(figsize=(20, 10))
```

```
plt.subplot(231)
```

```
sns.countplot(x='workclass',
```

```
hue='income',
```

```
data = df,
```

```
palette="BuPu")
```

```
plt.xticks(rotation=90)
```

```
plt.subplot(232)
```

```
sns.countplot(x='marital.status',
```

```
hue='income',
```

```
data = df,
```

```
palette="deep")
```

```
plt.xticks(rotation=90)
```

```
plt.subplot(233)
```

```
sns.countplot(x='education',
```

```
hue='income',
```

```
data = df,
```

```
palette = "autumn")
```

```
plt.xticks(rotation=90)
```

```
plt.subplot(234)
```

```
sns.countplot(x='relationship',
```

```
hue='income',
```

```
data = df,
```

```
palette = "inferno")
```

```
plt.xticks(rotation=90)
```

```
plt.subplot(235)
```

```
sns.countplot(x='sex',
```

```
hue='income',
```

```
data = df,
```

```
palette = "coolwarm")
```

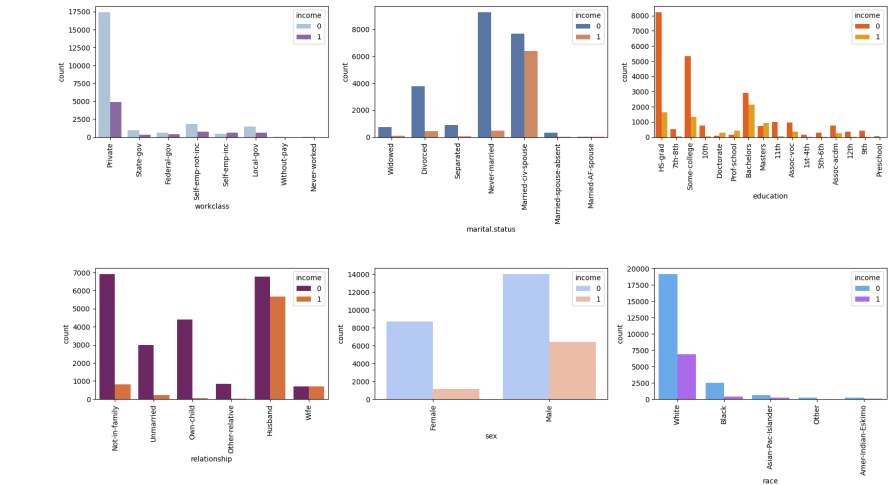
```
plt.xticks(rotation=90)
```

```
plt.subplot(236)
```

```
sns.countplot(x='race',
```

```
hue='income',
data = df,
palette = "cool1")
plt.xticks(rotation=90)
plt.subplots_adjust(hspace=1)
plt.show()
```

<ipython-input-16-e1b6d1f0108f>:3: MatplotlibDeprecationWarning: Auto-removal of o
plt.subplot(231)



```
df1 = df.copy()
categorical_features = list(df1.select_dtypes(include=['object']).columns)
print(categorical_features)
df1
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation
1	82	Private	132870	HS-grad	9	Widowed	Exec manageri
3	54	Private	140359	7th-8th	4	Divorced	Machine op-insp
4	41	Private	264663	Some-college	10	Separated	Pro specialt
5	34	Private	216864	HS-grad	9	Divorced	Other servic
6	38	Private	150601	10th	6	Separated	Adm clerics
...
32556	22	Private	310152	Some-college	10	Never-married	Protective ser
32557	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech supp

```

from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
for feat in categorical_features:
    df1[feat] = le.fit_transform(df1[feat].astype(str))
df1

X = df1.drop(columns = ['income'])
y = df1['income'].values
# Splitting the data set into train and test set
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 0)
print ("Train set size: ", X_train.shape)
print ("Test set size: ", X_test.shape)

from sklearn.ensemble import AdaBoostClassifier
# Train Adaboost Classifier
abc = AdaBoostClassifier(n_estimators = 300, learning_rate=1)
abc_model = abc.fit(X_train, y_train)
#Prediction
y_pred_abc = abc_model.predict(X_test)


Train set size: (21118, 14)
Test set size: (9051, 14)

print("Accuracy: ", accuracy_score(y_test, y_pred_abc))
print("F1 score :", f1_score(y_test, y_pred_abc, average='binary'))
print("Precision : ", precision_score(y_test, y_pred_abc))

Accuracy: 0.8637719588995691
F1 score : 0.7008007765105557
Precision : 0.7921009325287987

cm = confusion_matrix(y_test, y_pred_abc)
plt.figure(figsize=(5,5))
sns.heatmap(cm, annot=True, fmt=".3f", linewidths=.5, square = True, cmap ="coolwarm");
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
plt.title('Confusion Matrix - score:' + str(round(accuracy_score(y_test,y_pred_abc), 2)), size = 15);
plt.show()
print(classification_report(y_test, y_pred_abc))

```



```

from sklearn.ensemble import GradientBoostingClassifier
#Training the model with gradient boosting
gbc = GradientBoostingClassifier(
    learning_rate = 0.1,
    n_estimators = 500,
    max_depth = 5,
    subsample = 0.9,
    min_samples_split = 100,
    max_features='sqrt',
    random_state=10)
gbc.fit(X_train,y_train)
# Predictions
y_pred_gbc = gbc.predict(X_test)
print("Accuracy : ",accuracy_score(y_test, y_pred_gbc))
print("F1 score : ", f1_score(y_test, y_pred_gbc, average = 'binary'))
print("Precision : ", precision_score(y_test, y_pred_gbc))

```

```

Accuracy : 0.8689647552756602
F1 score : 0.7218574108818011
Precision : 0.7828077314343845

```

```

rms = np.sqrt(mean_squared_error(y_test, y_pred_gbc))
print("RMSE for gradient boost: ", rms)

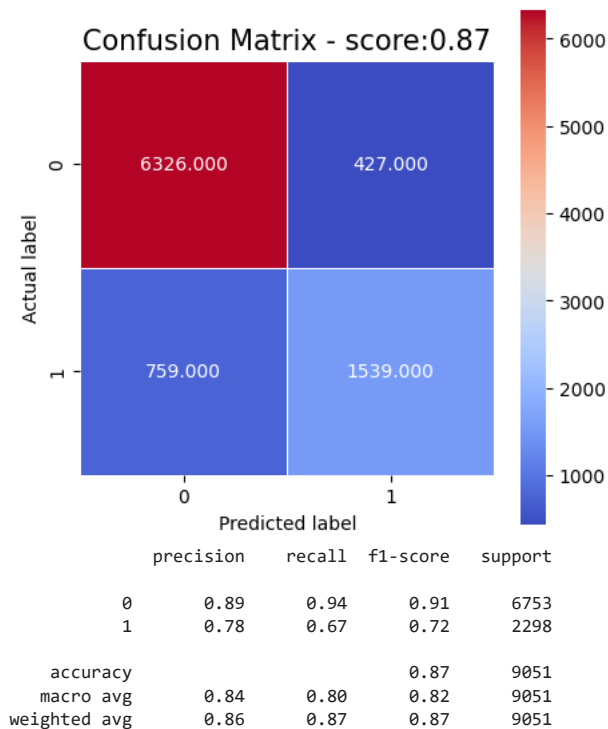
```

```

cm = confusion_matrix(y_test, y_pred_gbc)
plt.figure(figsize=(5,5))
sns.heatmap(cm, annot = True, fmt=".3f", linewidths = 0.5, square = True, cmap= "coolwarm");
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
plt.title('Confusion Matrix - score:' + str(round(accuracy_score(y_test,y_pred_gbc),2)), size = 15);
plt.show()
print(classification_report(y_test, y_pred_gbc))

```

RMSE for gradient boost: 0.3619879068758235



```

import xgboost as xgb
from xgboost import XGBClassifier
#Training the model with gradient boosting
xgboost = XGBClassifier(learning_rate=0.01,
    colsample_bytree = 0.4,
    n_estimators=1000,
    max_depth=20,
    gamma=1)
xgboost_model = xgboost.fit(X_train, y_train)


```

```
# Predictions
y_pred_xgboost = xgboost_model.predict(X_test)
print("Accuracy : ",accuracy_score(y_test, y_pred_xgboost))
print("F1 score : ", f1_score(y_test, y_pred_xgboost, average = 'binary'))
print("Precision : ", precision_score(y_test, y_pred_xgboost))

rms = np.sqrt(mean_squared_error(y_test, y_pred_xgboost))
print("RMSE for xgboost: ", rms)

cm = confusion_matrix(y_test, y_pred_xgboost)
plt.figure(figsize=(5,5))
sns.heatmap(cm, annot=True, fmt=".3f", linewidths=.5, square = True, cmap ="coolwarm");
plt.ylabel('Actual label');
plt.xlabel('Predicted label');
plt.title('Confusion Matrix - score:'+str(round(accuracy_score(y_test,y_pred_xgboost),2)), size = 15);
plt.show()
print(classification_report(y_test,y_pred_xgboost))
```

```
Accuracy : 0.8655397193680257
F1 score : 0.7090604829070045
Precision : 0.786737400530504
RMSE for xgboost: 0.3666882608319693
```



```
from sklearn.metrics import roc_curve
fpr, tpr, thresholds = roc_curve(y_test, xgboost.predict_proba(X_test)[:,:1])
plt.figure(figsize = (10,5))
plt.plot([0,1],[0,1], 'k--')
plt.plot(fpr, tpr)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC CURVE for Xgboost')
plt.show()
```

